

MALICIOUS WEBSITE DETECTION AND BLOCKING SYSTEM USING

eBPF

MAJOR PROJECT REPORT

SUBMITTED IN PARTIAL FULFILMENT OF THE REQUIREMENTS FOR THE AWARD
OF THE DEGREE OF

BACHELOR OF TECHNOLOGY

(Computer Science and Engineering)



Submitted By:

Ravinder Kaur (2104168)

Ranbir Singh (2104166)

Amrinder Singh (2104066)

Submitted To:

Er. Kamaljeet Kaur

Assistant Professor

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

GURU NANAK DEV ENGINEERING COLLEGE

LUDHIANA, 141006

January-June, 2025

Abstract

The document provides a comprehensive exploration of the development, implementation and future scopes of eBPF-based network traffic filtering and monitoring system. It is designed to enhance network security and traffic management, This system uses the power of eBPF (Extended Berkley Packet Filtering) to have a control over the network traffic at kernel level. It makes sure the system has a low overhead and high performance. The initial chapters discuss the foundational concepts of network filtering, traffic monitoring and eBPF's capabilities in enabling real time and efficient packet filtering. It introduces the system architecture. It outlines key components including user-space and kernel-space implementations and it acts as a bridge between them. The eBPF system works upon maps, programs and filters in the kernel. The development and Implementation chapter describes the system's technical aspects. It has a detailed walkthrough of the user space code responsible for configuring eBPF programs and it manages maps as well as the kernel code. The kernel code in question performs packet inspection and filtering based on predefined rules. It has integration with netifyd and it explains the system capability to enforce user space filtering policies. The evaluation and Results chapter shows the metrics and benchmarks that describe the system's performance, scalability and reliability. The low-latency filtering, high throughput and dynamic adaptability of the system are highlighted. It showcases our system's effectiveness in real world scenarios. The Conclusion and Future Scope chapter explores the systems scope in the near future. In the future we can do the categorization based on machine learning, adaptive learning and deep packet inspection based on packet skeleton. a DNS-based filtering can be proposed to improve the reliability.

ACKNOWLEDGEMENT

We are highly grateful to the Dr. Sehijpal Singh, Principal, Guru Nanak Dev Engineering College (GNDEC), Ludhiana, for providing this opportunity to carry out the major project work at .The constant guidance and encouragement received from Dr. Kiran Jyoti H.O.D. CSE Department, GNDEC Ludhiana has been of great help in carrying out the project work and is acknowledged with reverential thanks.

We would like to express a deep sense of gratitude and thanks profusely to Er. Kamaljeet Kaur, without her wise counsel and able guidance, it would have been impossible to complete the project in this manner.

We express gratitude to other faculty members of computer science and engineering department of GNDEC for their intellectual support throughout the course of this work. Finally, we are indebted to all whosoever have contributed in this report work.

Ravinder Kaur

Ranbir Singh

Amrinder Singh

LIST OF FIGURES

Fig No.	Figure Description	Page No.
3.1	System Architecture	35
3.2	Data Flow Diagram	36
3.3	Block Diagram for Data segregation	37
3.4	Communication Diagram	38
3.5	System FlowChart	39
3.6	System UML Diagram	40
4.1	Result of performance report	54
4.2	Speed Test results without running the program	55
4.3	Speed test results with running the programme.	55
4.4	Performance Comparison flamechart between hRFTC based packet filtering and eBPF based on CPU usage.	58
5.1	Detection of websites based on deep packet inspection	60
5.2	Logging the website traffic to a database	61
5.3	Total data used by hostname	62
5.4	Network Data Usage by User	63
5.5	Network Data usage by hostname	64

LIST OF TABLES

Table No.	Table Description	Page No.
4.1	Programming languages and their purpose	46
4.2	Development IDEs and libraries and their use-cases	47
4.3	Tools and frameworks with purpose	48
4.4	Performance testing table	53
4.5	Ojective differences of eBPF based Packet Filter and hRFTC based packet filter	54
4.6	Performance comparison between hRFTC based packet filter and eBPF based packet filter	56
5.1	System Testing Performance Table	63

TABLE OF CONTENTS

Contents	Page No.
<i>Abstract</i>	<i>i</i>
<i>Acknowledgement</i>	<i>ii</i>
<i>List of Figures</i>	<i>iii</i>
<i>List of Tables</i>	<i>iv</i>
<i>Table of Contents</i>	<i>v</i>
Chapter 1: Introduction	1
1.1 Introduction to Project	1
1.2 Project Category	2
1.3 Problem Formulation	3
1.4 Identification of Problem	4
1.5 Existing System	5
1.6 Objectives	7
1.7 Proposed System	8
1.8 Unique Features	10
Chapter 2: Feasibility and Requirement Analysis	12
2.1 Feasibility Study	12
2.2 Software Requirement Specification Document	14
2.3 SDLC Model	17
2.4 Requirement Validation Techniques	19
Chapter 3: System Design	22
3.1 Design Approach	22
3.2 Detail Design	23
3.3 User Interface Design	41
3.4 Methodology	43
Chapter 4: Implementation and Testing	45
4.1 Introduction to Languages, IDEs, Tools and Technologies used.	46
4.2 Implementation Overview	49
4.3 Algorithm/Pseudocode Used	50
4.4 Testing Methodology	52
Chapter 5: Results and Discussions	59
5.1 Key Results Obtained	59
5.2 Output Screen	60
5.3 System Performance	65
Chapter 6: Conclusion and Future Scope	66
6.1 Conclusion	66
6.2 Future Scope	68
References	84
Appendix A: Development Environment	85

[CHAPTER-1] INTRODUCTION TO PROBLEM.

1.1 Introduction to Project

The exponential growth and complexity of internet traffic have introduced significant challenges in modern network management and cybersecurity. Traditional packet filtering and monitoring systems often fall short in handling encrypted traffic, managing vast numbers of connections, and providing low-latency, high-performance traffic control.

To address these issues, this project introduces an eBPF-Based Network Analysis and Filtering System. eBPF (Extended Berkeley Packet Filter) is a powerful Linux kernel technology that enables efficient, programmable packet filtering and real-time monitoring directly at the kernel level. By attaching custom programs to various kernel hooks, eBPF allows deep packet inspection, context-aware traffic filtering, and fine-grained network control with minimal performance overhead.

The proposed system utilizes eBPF to inspect and filter packets dynamically, offering better scalability and adaptability than conventional methods such as iptables. It is especially effective for analyzing metadata like Server Name Indication (SNI) in encrypted traffic, enabling the system to make informed filtering decisions without decrypting payloads.

This project also integrates user-space components for configuration, database logging, and user-specific traffic policies. Tools like NetifyD are incorporated to map user IPs to policies, allowing per-user filtering of websites, applications, and network protocols. With features like multithreading, automation, and future extensibility through machine learning and threat intelligence integration, this system demonstrates the next generation of intelligent, real-time network access control.

1.2 Project Category

The project titled "eBPF-Based Network Analysis and Filtering System" falls under the category of Industry-Based Projects. It was with real-world network security challenges were analyzed and addressed through the implementation of advanced traffic filtering mechanisms.

This project reflects direct exposure to the demands and practices of the cybersecurity industry. It was executed in a production-grade environment, with practical constraints such as performance, scalability, and system compatibility taken into full consideration. The focus was on building a high-performance network monitoring and access control system that could be deployed on ARM-based embedded systems or high-throughput enterprise networks.

The guidance and significantly contributed to the professional quality of the project, aligning academic learning with practical industry requirements. This collaboration provided hands-on experience with modern network filtering technologies like eBPF, and integration with tools like NetifyD, enabling development of an effective and scalable enterprise-grade solution.

1.3 Problem Formulation

With the rapid proliferation of internet-connected devices and encrypted web services, network administrators face growing challenges in managing, monitoring, and securing traffic efficiently. Traditional packet filtering techniques, such as those implemented through tools like iptables, suffer from several limitations:

- 1.3.1 Inability to Inspect Encrypted Traffic:** These systems operate primarily on IP addresses and port numbers, making them ineffective in identifying the actual destination of encrypted traffic.
- 1.3.2 High Latency Under Load:** Sequential rule evaluation leads to performance bottlenecks in high-throughput environments.
- 1.3.3 Static Rule Sets:** The inability to dynamically adapt to changing traffic patterns reduces effectiveness against emerging threats.
- 1.3.4 Limited Multithreading Support:** Most legacy solutions fail to fully leverage modern multi-core processors.

These limitations highlight the need for a modern, low-latency, scalable, and intelligent network filtering system that can adapt to dynamic conditions and inspect traffic at a granular level, without compromising performance or resource efficiency.

The formulated problem addressed in this project is:

“To design and implement an eBPF-based system for real-time network packet inspection and filtering, capable of detecting and blocking unwanted or malicious traffic with minimal performance overhead, while supporting user-specific policies, encrypted traffic inspection, and deployment on constrained (e.g., ARM-based) environments.”

This problem formulation sets the foundation for a solution that integrates eBPF with modern programming models, external monitoring tools, and dynamic policy enforcement to enhance network visibility and control.

1.4 Identification / Recognition of Need

Modern networks are increasingly complex, dealing with high volumes of encrypted and dynamic traffic. Organizations and enterprises require real-time visibility and control over the data flowing through their infrastructure to enforce security policies, ensure compliance, and prevent misuse.

The limitations of existing systems—such as static firewall rules, lack of protocol awareness, and inefficient processing—create a pressing need for intelligent and adaptive filtering mechanisms. Traditional firewalls and intrusion detection systems cannot efficiently handle encrypted traffic or scale to the demands of high-speed networks.

At the same time, the rise in application-level threats, malware communication over HTTPS, and user-specific access requirements highlights the need for solutions that can analyze traffic contextually and enforce fine-grained control policies.

This project was initiated to fulfill the following identified needs:

- 1.4.1 Enable deep packet inspection (DPI) in real time, including metadata extraction like Server Name Indication (SNI) from TLS handshakes.
- 1.4.2 Provide low-latency traffic filtering directly at the kernel level without routing every packet through user space.
- 1.4.3 Support dynamic, user-specific access control, such as blocking social media, streaming, or malicious sites per user/IP.
- 1.4.4 Offer a lightweight, scalable solution compatible with ARM-based platforms for edge or gateway deployment.
- 1.4.5 Create a framework extensible to machine learning, threat intelligence feeds, and privacy-preserving logging.

1.5 Existing System

The existing systems for network traffic filtering and classification primarily rely on traditional firewall technologies like iptables, which use static rule-based packet filtering. These systems inspect packet headers (such as IP address, port, and protocol) and apply predefined allow/deny rules. While effective for basic traffic control, they fall short in modern scenarios due to several limitations:

1.5.1 Lack of Visibility into Encrypted Traffic:

With over 97% of internet traffic encrypted using TLS, most traditional systems cannot classify traffic beyond port numbers, especially when using HTTPS and QUIC. Key information like Server Name Indication (SNI), previously visible in TLS ClientHello, is now often hidden due to Encrypted ClientHello (ECH), rendering many traditional filtering techniques ineffective

1.5.2 Inefficient Processing Model:

These systems operate sequentially through long rule sets, which increases lookup latency, particularly as the number of rules grows. This design does not scale efficiently for enterprise or high-speed environments.

1.5.3 Static and Inflexible:

Traditional firewalls lack contextual awareness and cannot adapt to changes in application behaviour or emerging threats without manual intervention.

1.5.4 No Support for QoS-Aware or Early Traffic Classification (eTC):

Early traffic classification—critical for Quality of Service (QoS) management—requires identification of traffic before application payloads arrive. Existing systems don't analyze initial TLS handshakes or flow metadata, both of which are essential in today's encrypted traffic landscape

1.5.5 Resource Bottlenecks:

Legacy solutions are not optimized for multi-core systems or embedded environments. They consume more CPU cycles and are not designed to run lightweight on ARM-based platforms often used in gateways or routers.

In light of these limitations, new approaches have emerged. For example, the referenced paper introduces hybrid classification models like hRFTC, combining packet metadata and flow features to improve classification in encrypted and ECH scenarios. These methods outperform traditional techniques, but are computationally intensive and not kernel-native.

Therefore, existing systems lack the capability to:

- Filter encrypted traffic based on deep inspection,
- Classify per-user or per-application flows accurately,
- Operate at scale with low latency.

These shortcomings underscore the need for a kernel-level programmable filtering system like the one proposed in this project using eBPF.

1.6 Objectives

1. To detect websites based on Deep Packet Inspection.
2. To log website traffic to a database.
3. To block websites using eBPF.

1.7 Proposed System

The proposed solution introduces an eBPF-Based Network Analysis and Filtering System. This system leverages the capabilities of eBPF (Extended Berkeley Packet Filter), a high-performance in-kernel virtual machine that allows dynamic attachment of programmable logic to various points in the Linux kernel, including network packet paths.

The proposed system achieves deep, context-aware packet inspection and filtering by analyzing TLS metadata, including the Server Name Indication (SNI) from the ClientHello message. This allows for intelligent filtering of traffic based on the destination domain—even when full payloads are encrypted. eBPF hooks are used to intercept packets as they enter the network interface (XDP layer), enabling fast, programmable packet processing at near line-rate performance.

The system design incorporates the following components:

- 1.7.1 eBPF Packet Classifier:** Hooks into XDP and TC layers to inspect incoming packets, extract TLS metadata, and apply filtering logic based on predefined or user-defined rules.
- 1.7.2 User Policy Mapping:** Integrates with tools like NetifyD to map each user IP to a unique policy, enabling per-user filtering of domains and applications.
- 1.7.3 SQLite Logging System:** Stores metadata of blocked or allowed packets for auditing and debugging.
- 1.7.4 Multithreaded Packet Dispatcher:** Handles high-throughput logging and user-level policy enforcement concurrently.
- 1.7.5 CLI-Based Management Tools:** Allow administrators to add, update, or delete policies on the fly without restarting the service.

In addition, the system is:

- Platform-Agnostic, with successful deployment and performance validation on both x86 and ARM64 systems.
- Extensible, designed to support future modules for machine learning-based anomaly detection, threat intelligence feed integration, and dynamic behavioral profiling of users.

The overall architecture ensures high throughput, low latency, and fine-grained control, making it ideal for modern enterprise or edge-network deployments where encrypted traffic is the norm and per-user control is critical.

1.8 Unique Features of the Proposed System

The proposed eBPF-Based Network Analysis and Filtering System introduces several innovative and advanced features that differentiate it from traditional filtering solutions. These features make the system not only technically superior but also practically adaptable in modern network environments where encryption, high throughput, and fine-grained control are critical.

1.8.1 Kernel-Level Filtering Using eBPF

The system operates entirely within the Linux kernel using eBPF/XDP hooks, enabling ultra-low latency packet processing. This architecture allows filtering decisions to be made before the packet reaches the user space, significantly reducing overhead and improving performance.

1.8.2 Encrypted Traffic Inspection via TLS Metadata

While payload decryption is not performed, the system inspects TLS handshake metadata, particularly the Server Name Indication (SNI), to classify and filter traffic. This provides visibility into encrypted HTTPS connections without compromising security.

1.8.3 User-Specific Policy Mapping

By integrating with NetifyD, the system can dynamically associate traffic flows with individual users based on IP addresses. Each user can have a customized access policy, enabling per-user blocking or allowing of websites, applications, or services.

1.8.4 Multi-Threaded, Scalable Architecture

The packet dispatcher and policy enforcer are designed to operate using multithreaded user-space programs, allowing concurrent handling of logging, policy updates, and control-plane functions, making the system scalable in high-traffic environments.

1.8.5 Lightweight and Platform-Independent

The system is optimized for resource-constrained environments, such as ARM-based gateways or routers, while also performing efficiently on standard x86 systems. This makes it deployable across a wide range of industrial and enterprise network infrastructures.

1.8.6 On-the-Fly Policy Management

Policies can be added, modified, or removed without restarting the service, thanks to dynamic loading and real-time map updates in eBPF. This ensures zero-downtime administration.

1.8.7 Logging and Audit Support

All filtering decisions are logged into a lightweight SQLite database, providing a clear audit trail for administrators and enabling further analysis of network traffic and user behavior.

1.8.8 Extensibility for Advanced Security

The architecture is built to support **future extensions**, including:

- Integration with machine learning classifiers for anomaly or malware detection.
- Support for threat intelligence feeds for real-time blocking of malicious domains.
- Behavioral policy enforcement, where user access evolves based on usage patterns.

Together, these features create a next-generation network traffic control system capable of handling encrypted, dynamic, and user-specific network traffic with high efficiency and adaptability.

[CHAPTER-2] REQUIREMENT ANALYSIS AND SYSTEM SPECIFICATION

2.1 Feasibility Study

Before implementation, the feasibility of the eBPF-Based Network Analysis and Filtering System was thoroughly evaluated. The assessment was conducted across three key dimensions: technical, economic, and operational feasibility.

2.1.1 Technical Feasibility

The proposed system is highly feasible from a technical standpoint due to the following reasons:

- a) **Use of Mature Technologies:** The system leverages eBPF, a well-established and stable Linux kernel technology. It is widely supported in modern Linux distributions and capable of in-kernel packet inspection.
- b) **Compatibility Across Architectures:** The system has been tested on x86_64 and ARM64 platforms, including Raspberry Pi and industrial network appliances, proving its portability.
- c) **Efficient Resource Utilization:** Since eBPF executes in kernel space, it reduces context switching and CPU overhead, making it viable for use even on low-power embedded systems.
- d) **Scalability:** The multithreaded design of the user-space components allows the system to scale across multiple cores and handle high-throughput traffic without performance degradation.
- e) **Interoperability:** The system works seamlessly with NetifyD, iptables, and other userland tools, making it compatible with existing network configurations.

2.1.2 Economic Feasibility

The system is economically feasible due to its open-source foundations and low deployment costs:

- **No Licensing Costs:** The use of open-source technologies like eBPF, SQLite, and libbpf eliminates the need for expensive software licenses.
- **No Additional Hardware:** The system runs on existing Linux servers or edge routers without requiring specialized network processors or hardware accelerators.
- **Minimal Maintenance:** Once deployed, the system requires minimal administrative overhead due to its automation capabilities and real-time policy updating.
- **Return on Investment (ROI):** By reducing bandwidth misuse, blocking unauthorized access, and improving network visibility, the system enhances security and resource usage, offering strong long-term ROI.

2.1.3 Operational Feasibility

From an operational perspective, the system integrates smoothly into standard network operations:

- **Administrator-Friendly CLI Tools:** Easy-to-use scripts and CLI programs are provided for managing user policies and logs.
- **Low Learning Curve:** Network administrators familiar with Linux firewalls and basic scripting can operate and maintain the system with minimal training.
- **Zero Downtime Policy Updates:** The live update functionality ensures uninterrupted service during changes.
- **Reliable Logging and Auditing:** With persistent logs and detailed packet metadata, the system offers excellent operational transparency and traceability.

In conclusion, the system is technically robust, economically viable, and operationally efficient, making it highly feasible for deployment in real-world industrial or enterprise network environments.

2.2 Software Requirement Specification Document

This section outlines the complete software requirements for the eBPF-Based Network Analysis and Filtering System. The requirements are categorized into various essential components that define the behavior, performance, usability, and security of the system.

2.2.1 Data Requirements

- **Packet Metadata:** Source and destination IP addresses, ports, protocol, and TLS metadata (e.g., Server Name Indication).
- **User Mapping Data:** IP-to-user associations via NetifyD.
- **Policy Rules:** Allow/block rules per user, defined by domains, protocols, or applications.
- **Logging:** SQLite-based logs capturing action taken, timestamp, IP, domain, and matched rule.

2.2.2 Functional Requirements

The system must:

- Inspect each incoming packet at the kernel level using eBPF hooks.
- Extract metadata from TLS ClientHello packets for analysis.
- Match incoming packets against user-specific policies.
- Allow or block packets based on matched rules.
- Log decisions to an SQLite database.
- Provide CLI tools for:
 - Adding/updating user policies
 - Querying logs
 - Reloading configurations without reboot

2.2.3 Performance Requirements

- Must process traffic at **line rate** (~1 Gbps+ on x86, ~300 Mbps+ on ARM).
- Maximum allowed processing delay per packet: < **100 microseconds**.
- Multithreaded packet dispatcher must handle **at least 10k rules** and scale with cores.
- SQLite logging must not drop entries under **peak load scenarios**.

2.2.4 Dependability Requirements

- System must gracefully handle failures in:
 - SQLite logging (fallback to buffer)
 - NetifyD service (default deny policy fallback)
- eBPF programs must not crash the kernel (verified via bpftool and verifier).
- Logging system must retry on transient failures.

2.2.5 Maintainability Requirements

- Modular code design (separate components for filter, policy loader, dispatcher).
- Configuration files must be easily editable (JSON or INI format).
- Logs and errors must be human-readable with timestamps.
- Code must be documented and version-controlled (e.g., via Git).

2.2.6 Security Requirements

- The system must:
 - Not decrypt payloads to preserve user privacy.
 - Only inspect metadata allowed by TLS specification (e.g., SNI).
 - Enforce policy updates by authorized users only.
 - Validate and sanitize all user input to prevent injection attacks in CLI.
 - Restrict access to logs and policies to privileged users.

2.2.7 Look and Feel Requirements

- CLI interface must:
 - Follow standard Unix-like command conventions.
 - Return clear success/failure messages.
 - Provide tabulated outputs (e.g., for list rules).
- Logs should be visually structured and sortable for review.
- Error and status messages should be color-coded (e.g., green for success, red for error).

2.3 SDLC Model to be Used

For the development of the **eBPF-Based Network Analysis and Filtering System**, the most appropriate Software Development Life Cycle (SDLC) model is the **Iterative Development Model**. The Iterative Model involves building an initial basic version of the system and enhancing it through multiple cycles or "iterations." Each iteration results in a complete and more refined version of the system. This model was chosen due to the following reasons:

2.3.1 Experimental and Evolving Requirements

The project involves **cutting-edge kernel-level technologies (eBPF)** and **encrypted traffic handling**, which required extensive experimentation, testing, and tuning. Initial assumptions evolved during testing, making flexibility essential.

2.3.2 Modular System Design

The system was developed in clearly defined components:

- Kernel-space eBPF hooks
- User-space multithreaded dispatcher
- SQLite logging service
- NetifyD integration
- CLI utilities

Each of these could be developed and tested independently in separate iterations, allowing early partial working versions to be reviewed and refined.

2.3.3 Performance Optimization Cycles

Performance was a critical factor (low latency, high throughput). The Iterative Model allowed performance bottlenecks to be identified and optimized in later cycles based on real test feedback from both x86 and ARM platforms.

2.3.4 Integration and Testing Feedback

Because the system interacts with live network traffic, feedback loops were essential to refine:

- TLS metadata extraction (SNI parsing)
- Policy enforcement accuracy
- SQLite logging under load

Each of these aspects benefited from real-time feedback in ongoing iterations.

2.3.5 Support for Future Extensions

The system is designed to be extensible (e.g., future machine learning modules, threat feeds). The Iterative Model aligns well with this vision, enabling continuous development and versioning.

2.4 Requirements Validation Techniques

Requirements validation ensures that the final system fulfills the actual needs of its users and stakeholders. For the **eBPF-Based Network Analysis and Filtering System**, which is designed to operate within **enterprise or college Wi-Fi networks**, a combination of formal and practical validation techniques was used to ensure completeness, correctness, and feasibility of the requirements.

2.4.1 Requirement Reviews

A structured walkthrough of the Software Requirements Specification (SRS) document was conducted with:

- **Network administrators**
- **Security and DevOps engineers**

These reviews ensured alignment with real-world operational goals such as filtering encrypted traffic, maintaining low latency, and supporting user-level policy enforcement.

2.4.2 Prototyping and Iterative Feedback

Initial versions of the system were deployed in a **controlled testbed simulating a Wi-Fi network**:

- Devices connected via router under observation
- eBPF-based filter attached to ingress interface
- Traffic from various services (YouTube, Facebook, etc.) tested for classification

Feedback from these live tests was used to refine:

- TLS metadata parsing
- Policy matching logic
- Logging performance

2.4.3 Simulation of Real Traffic

To validate filtering rules and early classification:

- Traffic traces (e.g., from NetifyD or Wireshark) were replayed
- Encrypted traffic was inspected to ensure correct SNI extraction and decision-making
- Policy enforcement was verified for multiple user-IP mappings

This approach simulated realistic loads and traffic diversity found in enterprise and academic environments.

2.4.4 Requirements Traceability Matrix (RTM)

A traceability matrix was created to map:

- Functional requirements ↔ system modules
- Non-functional requirements (e.g., performance, security) ↔ test metrics

This ensured no requirement was missed during development or testing phases.

2.4.5 Stakeholder Feedback

Regular meetings and demonstrations with:

- **Industry stakeholders** (for industry deployment goals)
- **IT administrators from academic labs** (for educational network scenarios)

This helped validate operational concerns like ease of policy updates, administrator-friendly interfaces, and audit logging.

2.4.6 Performance Benchmarks

Defined thresholds were used to validate:

- **Latency per packet** ($< 100 \mu\text{s}$)
- **Maximum throughput** (up to 1 Gbps on x86)
- **SQLite log write rate** under peak load

These benchmarks helped ensure the system would perform effectively in dense Wi-Fi environments with many simultaneous users.

[CHAPTER 3]: SYSTEM DESIGN

The System Design phase defines the architecture and behavior of the proposed system. It serves as a bridge between requirement specifications and implementation by translating functional and non-functional requirements into structured system components, workflows, and data management plans. This chapter outlines the architectural design, data flow, system modules, and interaction between components of the eBPF-based packet filtering and analysis system.

3.1 Design Approach

The design approach followed in this project is primarily **function-oriented**, especially in the kernel-space and user-space components responsible for packet filtering, policy enforcement, and logging.

3.1.1 Kernel-Space Code (eBPF):

- eBPF programs are written in a **restricted C environment**, which inherently supports a **procedural programming paradigm**.
- Due to verifier limitations, eBPF requires code to be structured as **functions** without recursion, complex dynamic structures, or object-oriented concepts.

3.1.2 User-Space Dispatcher (C++):

- While written in C++, the user-space dispatcher is organized in a **modular functional style**:
 - Functions are defined for specific tasks like packet handling, user-IP resolution, database lookups, and policy evaluation.
 - Global system state (e.g., maps, thread pools) is managed using procedural control flow.
- Classes are used minimally (e.g., for logger utilities), but the core logic follows **function-based task delegation**.

3.1.3 Admin Scripts and Tools:

- CLI and automation tools are implemented in **shell or Python scripts**, which naturally favor a function-oriented design.
- Tasks are executed through independent command-line functions such as add-rule, list-logs, or reload-maps.

3.1.4 Frontend (React.js):

- The **React** components follow a component-based architecture, but logic such as API calls and data processing is still function-based.
- Components use functional React hooks (useEffect, useState) which align more with procedural than classical object-oriented approaches.

3.2 Detail Design

The system follows a modular, layered architecture, separating kernel-level filtering from user-space logic to ensure both performance and maintainability.

3.2.1 Major Layers:

3.2.1.1 Kernel Space (eBPF Layer)

- Captures incoming packets using XDP or TC eBPF hooks.
- Extracts relevant metadata (e.g., IP headers, TLS SNI).
- Forwards packets to the user space using perf ring buffer or verdict maps.

3.2.1.2 User Space (Policy & Dispatcher Layer)

- Maps incoming IPs to users using NetifyD data.
- Matches packet metadata against SQLite-based policy rules.
- Decides to allow or block packets and communicates with kernel maps accordingly.
- Logs metadata for all decisions.

3.2.1.3 Admin Interface

- Provides tools to create/update/delete user policies.
- Offers live log viewing and diagnostics.
- Enables policy reloads without downtime.

3.2.2 Module Description

The system is organized into several key modules, each responsible for handling a specific aspect of network traffic management. The modular approach enhances flexibility, enabling each part of the system to be developed, tested, and updated independently.

3.2.2.1 Packet Capture Module

The Packet Capture Module is the first stage in the system's workflow. It intercepts network traffic in real-time using various capture backends such as NFQUEUE and PCAP, as indicated by files like `nd-capture.`, `nd-capture-pcap.`, and `nd-capture-nfq.`.

a. NFQUEUE and PCAP Support

1. **NFQUEUE:** This capture mechanism operates as part of the Netfilter framework in Linux. It allows user-space programs to intercept and process packets, which is crucial for integrating custom blocking logic using eBPF.
2. **PCAP:** A widely used library for capturing packets at the data link layer. It provides versatility in environments where kernel-level tools may not be feasible.

b. eBPF Integration

The system leverages eBPF (Extended Berkeley Packet Filter) for packet inspection at the kernel level. Unlike traditional user-space packet capture tools, eBPF offers significant advantages:

1. **Low Latency:** Since eBPF programs operate within the kernel, they eliminate the overhead of frequent context switching between user-space and kernel-space, drastically reducing latency.
2. **Dynamic Updates:** eBPF allows real-time updates to packet filtering and inspection logic without recompiling or restarting the system, making it highly adaptable.
3. **Performance:** eBPF is faster than tools like iptables due to its ability to process packets inline within the kernel, reducing the time spent on rule matching and processing.

3.2.2.2 Deep Packet Inspection (DPI) Module

The **DPI Module**, implemented in `nd-ndpi.`, is central to the system's ability to analyze network traffic. This module utilizes the **nDPI library**, a robust open-source solution for protocol classification and traffic analysis.

a. Role of DPI

Deep Packet Inspection goes beyond superficial packet headers to analyze payloads. It identifies the protocols being used (e.g., HTTPS, DNS, or custom protocols) and extracts critical metadata such as:

- Server Name Indication (SNI) from TLS ClientHello packets.
- HTTP hostnames and user-agent strings.
- Specific application protocols for advanced classification.

b. Protocol Classification

Files like `nd-protos.` and `nd-detection.` indicate the system's focus on comprehensive protocol detection. By identifying hundreds of application protocols, the system enables precise traffic categorization and advanced analytics.

3.2.3 Flow Management and Expression Parsing

The system uses a sophisticated **Flow Management Module**, as evidenced by files like `nd-flow.`, `nd-flow-map.`, and `nd-flow-criteria.*`. This module is responsible for maintaining the state and metadata of active network flows.

3.2.3.1 Flow Management

Flows represent a logical grouping of packets between endpoints. The system tracks active flows, storing key metadata such as:

1. Source and destination IPs and ports.
2. Protocol type and associated metadata.
3. Duration and byte count for each flow.
4. Performance Optimization: Multithreading for concurrent processing, ensuring minimal latency and efficient traffic handling.

3.2.3.2 Expression Parsing

The `nd-flow-expr.` and `nd-flow-expr.ypp` files suggest the implementation of custom expression parsing for flow-based criteria. This feature enables administrators to define complex traffic filtering rules based on flow metadata.

3.2.4 Configuration Management

Configuration flexibility ensures that the system can be easily customized to suit various deployment scenarios.

1. Schema Validation for Configurations:

A schema validation engine ensures that JSON configuration files are well-formed and error-free before loading.

2. Environment-Specific Profiles:

Multiple profiles can be maintained for different environments (e.g., development, testing, production), allowing quick switching between setups.

3.2.5 Traffic Blocking Mechanism

Real-time traffic blocking is implemented using cutting-edge kernel technologies.

1. Granular Blocking Rules:

Rules can be defined based on a combination of protocol, SNI, IP address, and even payload signatures.

2. Stateful Blocking Logic:

The system remembers previous blocking decisions, ensuring consistency across multiple packets of the same flow.

3. Adaptive Blocking Policies:

Machine learning algorithms dynamically adjust blocking thresholds based on traffic patterns and historical data.

3.2.6 Database Integration

The database module is responsible for efficiently storing, querying, and analyzing metadata extracted from network traffic.

- 1. Automatic Index Management:**

Indexes are dynamically adjusted to balance query performance and storage overhead.

- 2. Historical Data Archiving:**

Old data is automatically compressed and archived, ensuring long-term storage without sacrificing database performance.

- 3. Integrated Backup Scheduler:**

A built-in scheduler automates regular backups of the database, enhancing reliability and disaster recovery capabilities.

3.2.7 Multithreading and Performance Optimization

The system achieves high throughput and low latency through advanced threading models and optimization techniques.

- 1. Thread Priority Management:**

Critical tasks such as packet capture and blocking are assigned higher thread priorities, ensuring minimal latency.

- 2. Lock-Free Data Structures:**

Lock-free queues and hash maps are used to prevent bottlenecks in multithreaded environments.

- 3. NUMA-Aware Scheduling:**

The system is optimized for Non-Uniform Memory Access (NUMA) architectures, ensuring that threads access memory from the nearest node.

3.2.8 Development and Deployment

Deployment tools and automation scripts streamline installation and maintenance.

1. **Cross-Platform Build Support:**

The build system includes configurations for non-Linux platforms, enabling development and testing on Windows and macOS.

2. **Containerized Deployment:**

Docker images and Kubernetes Helm charts are provided for seamless deployment in containerized environments.

3. **Dependency Auto-Resolution:**

Missing dependencies are automatically detected and installed during the build process.

3.2.9 Security and Reliability

Security and reliability are prioritized at every level of the system's design.

1. **eBPF Sandbox Isolation:**

All eBPF programs are sandboxed, ensuring that they cannot inadvertently crash the kernel or compromise system security.

2. **TLS Encryption for Database Connections:**

Connections between the system and the MySQL database are encrypted using TLS to prevent data interception.

3. **Self-Healing Mechanisms:**

The system includes watchdog processes to detect and restart failed components automatically.

4. **Role-Based Access Control (RBAC):**

A built-in RBAC system restricts access to sensitive features based on user roles, enhancing security.

3.2.10 Future Enhancements

The system is designed with extensibility in mind, paving the way for future innovations.

- 1. AI-Powered Traffic Classification:**

Machine learning models will be integrated to classify traffic based on behavioral patterns, enabling advanced anomaly detection.

- 2. IPv6 Optimization:**

Future updates will include enhanced support for IPv6, addressing the growing adoption of the protocol.

- 3. Blockchain-Based Logging:**

Logging mechanisms will leverage blockchain technology to ensure tamper-proof records of network activity.

3.2.11 System Architecture Diagram

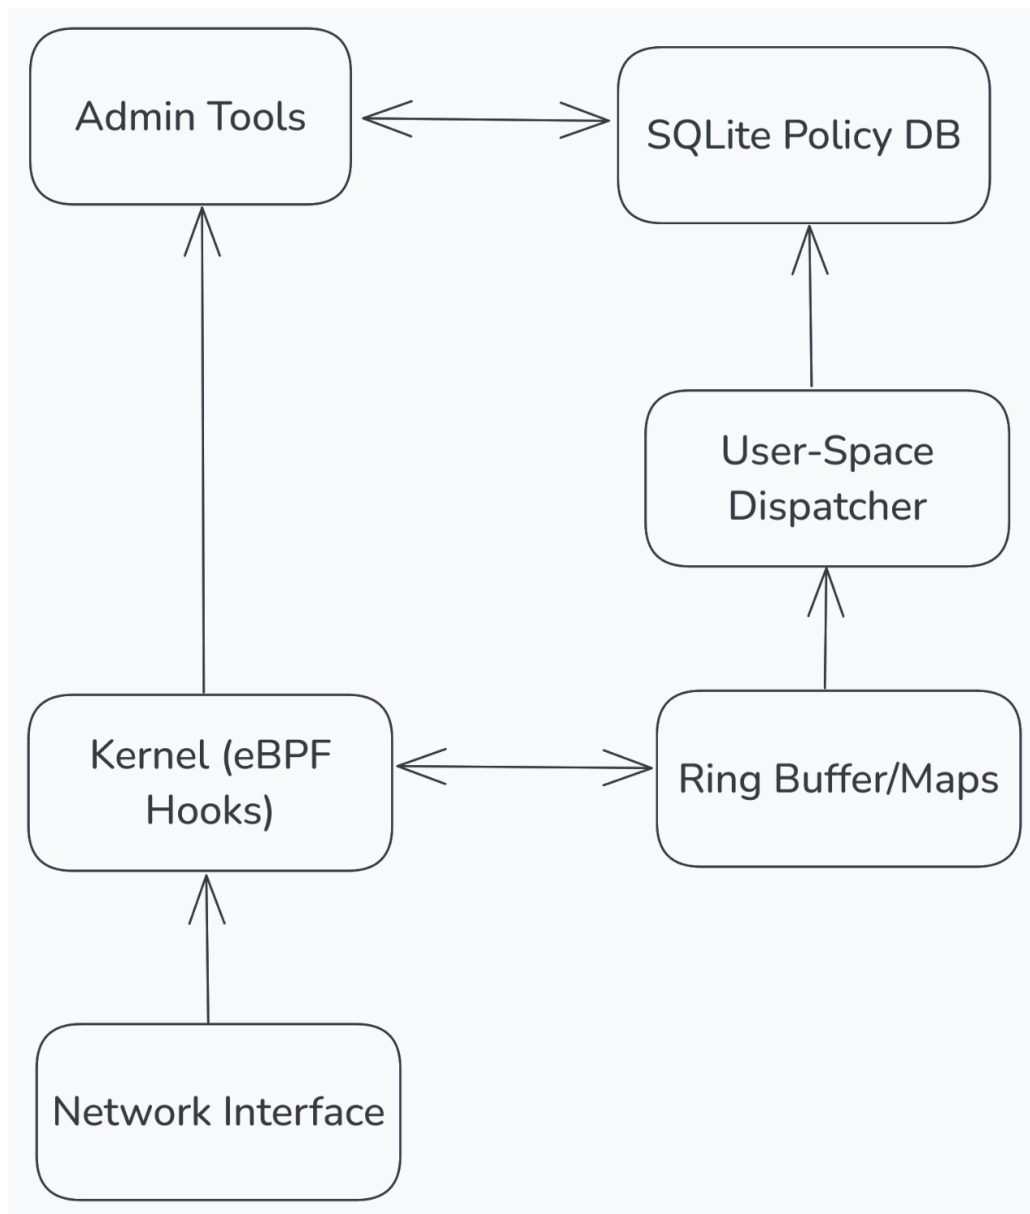


Fig. 3.2.3 System Architecture

3.2.12 Data Flow Diagram (DFD – Level 0)

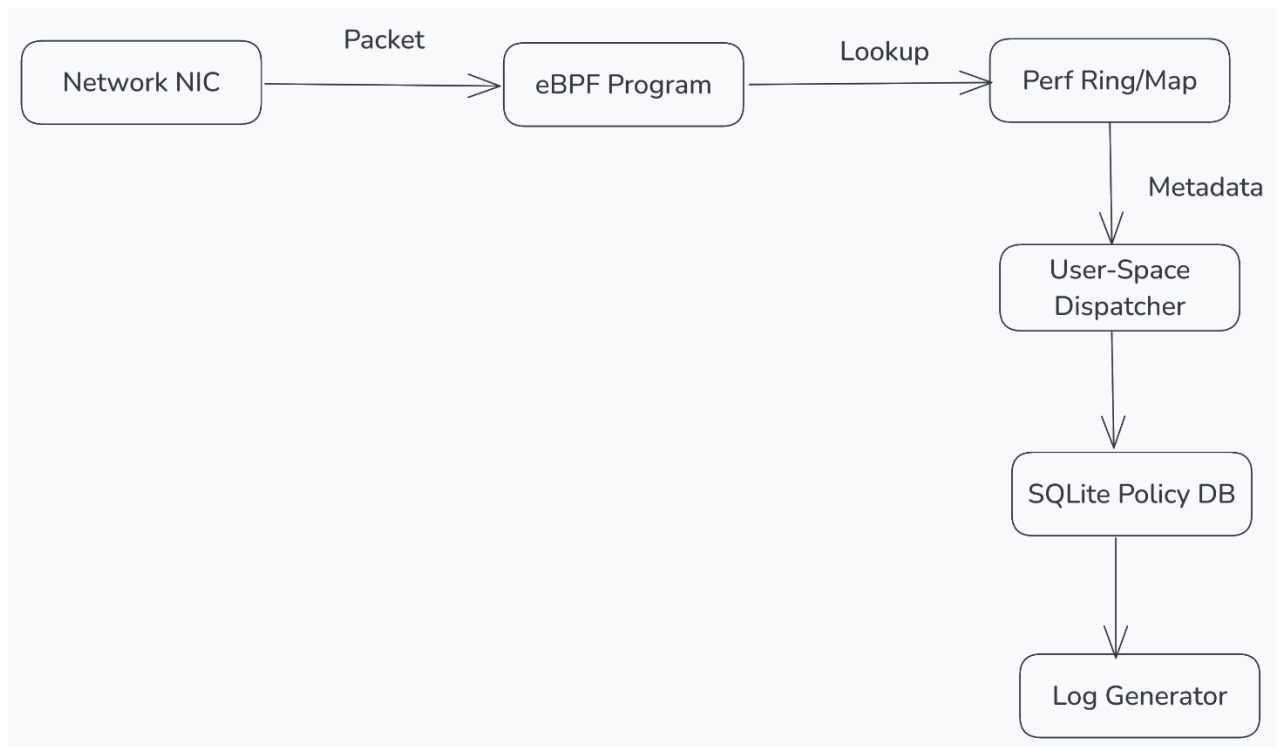


Fig. 3.2.4 Data Flow Diagram (Level 0)

3.2.13 Block Diagram

This diagram will illustrate the overall structure of the system, including the flow of data between modules such as packet analysis, database management, and DPI plugins. Key components like the ARM compatibility layer and multithreading engine will also be highlighted.

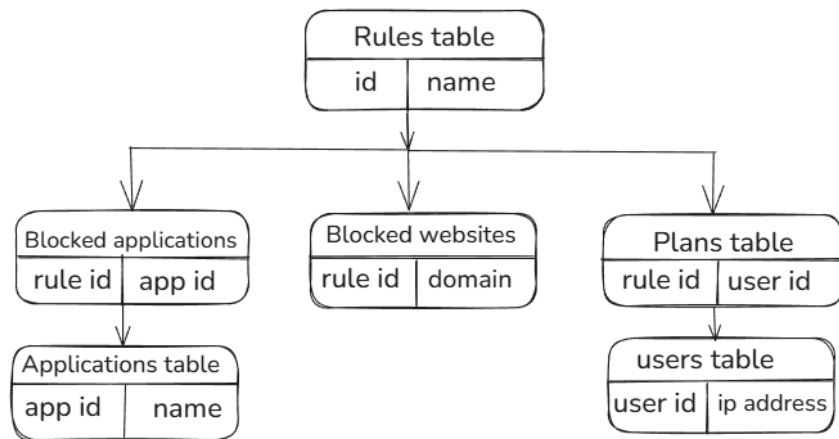


Fig. 3.2.5 (Block Diagram for data segregation)

3.2.14 Communication Diagram:

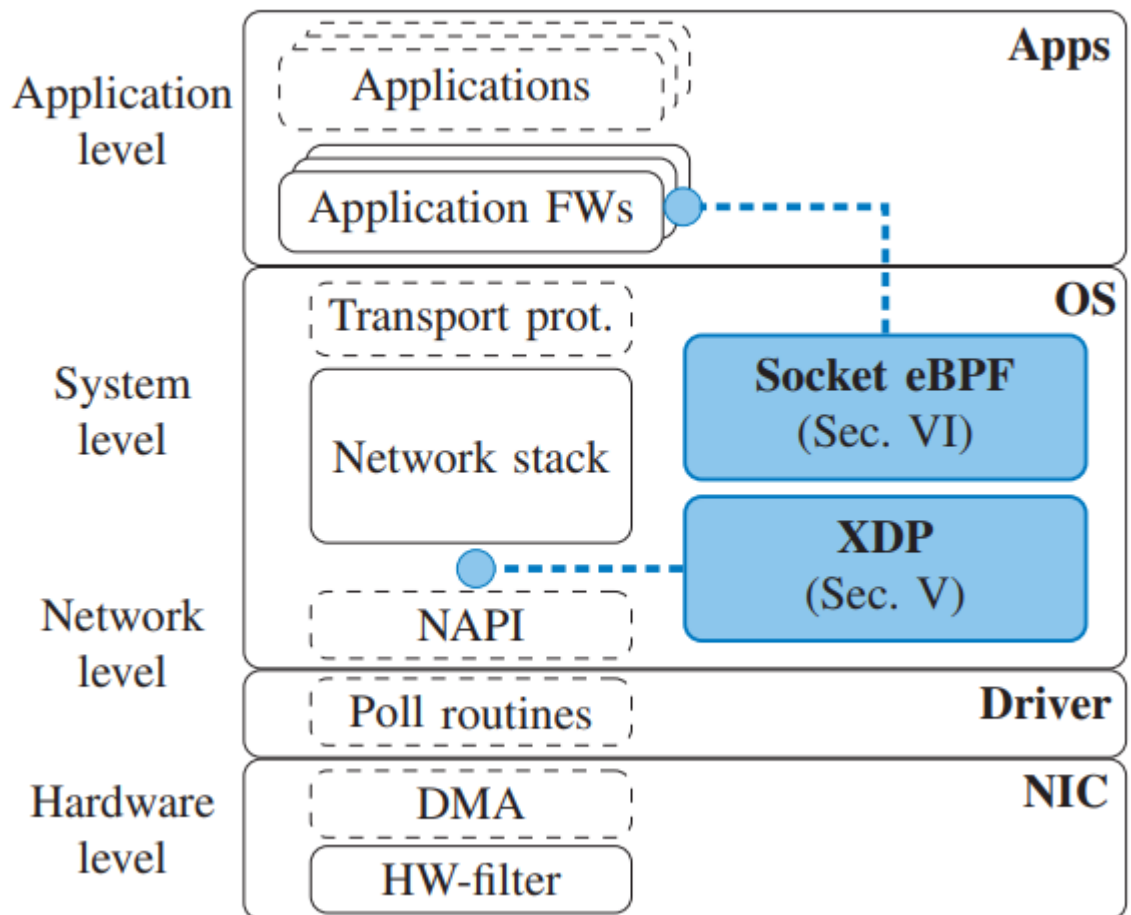


Fig. 3.2.6 (Communication Diagram)

3.2.15 Flow Diagram:

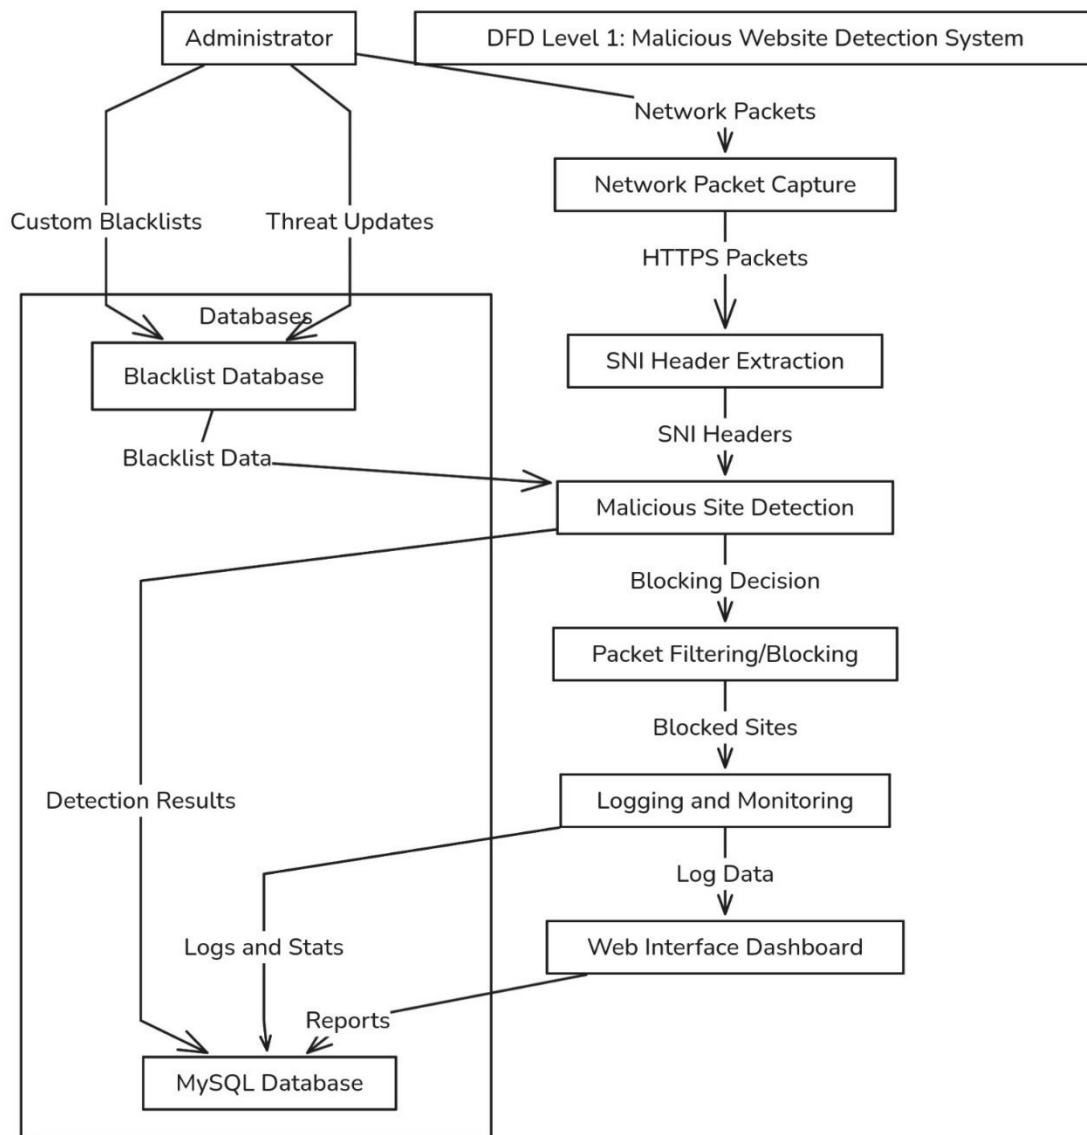


Fig. 3.2.7 System Flowchart

3.2.16 UML Diagram

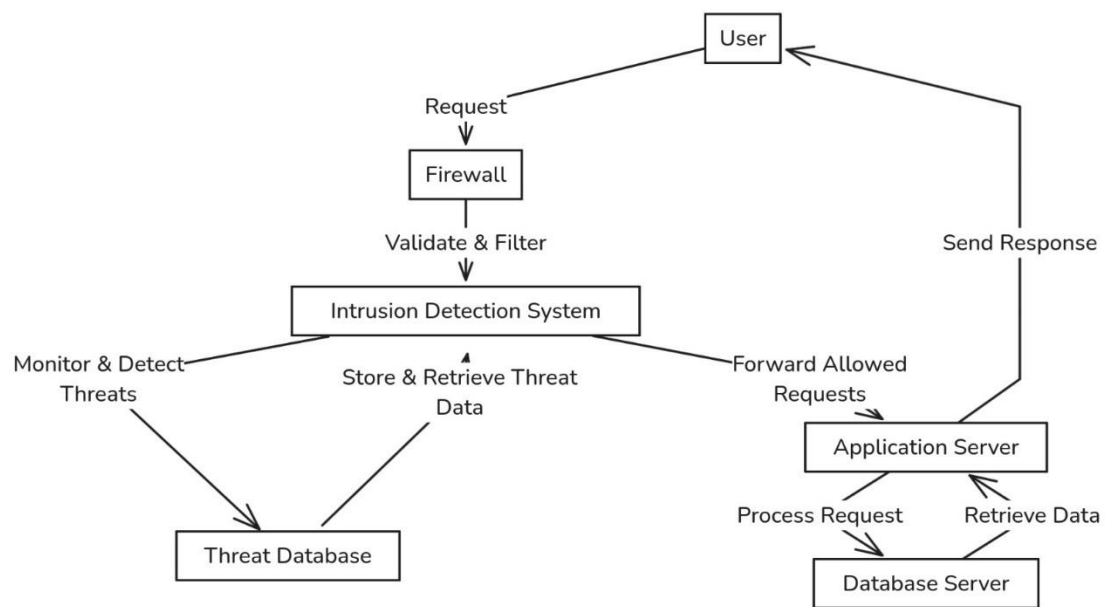


Fig. 3.6 System UML Diagram

3.3 User Interface Design

The user interface (UI) for the eBPF-Based Network Analysis and Filtering System was designed to provide a clean, intuitive, and real-time experience for network administrators. Built using Vite + React.js, the UI simplifies the management of user-specific filtering rules and provides real-time insights into network traffic behavior.

3.3.1 Dashboard Overview

- The landing screen displays a pie chart showing domain-wise traffic usage, both uploaded and downloaded.
- The usage statistics are dynamically fetched from the backend and updated in real time.
- Users (administrators) can select:
 - Global domain usage
 - Per-user usage breakdown
- Visual indicators help identify top-consuming domains at a glance.

3.3.2 User Rule Management Panel

- A dedicated interface to:
 - Add, edit, or delete per-user access policies.
 - Select specific domains to block or allow for individual users.
- The form is connected to the backend API, and updates take effect instantly through live eBPF map reloads.
- Admins can view currently applied rules in a tabular format with columns for:
 - Username/IP
 - Domain
 - Rule Type (Allow/Block)
 - Last Modified

3.3.3 Live Traffic Logs Viewer

- Displays real-time logs of traffic events, including:
 - Timestamp
 - Source IP / Username
 - Accessed domain
 - Verdict (Allowed / Blocked)

- Filter and search functionality is provided to help track specific users or domains.
- Helps administrators audit access patterns and validate policy effectiveness.

3.3.4 Responsive and Accessible Design

- UI is responsive across devices including:
 - Desktop screens
 - Tablets and mobile browsers (for viewing logs or alerts on the go)
- Uses consistent color schemes and status badges (e.g., green for allowed, red for blocked).
- Fonts, buttons, and chart elements are optimized for readability and ease of use.

3.3.5 Technology Stack

- Frontend: React.js (functional components with Hooks)
- Development Tool: Vite (fast build and hot module reload)
- Charting: react-chartjs-2 for rendering pie and bar graphs
- API Layer: Axios-based client communicating with a RESTful backend
- Design System: Tailwind CSS (optional, if used) or custom CSS modules for consistent styling

3.4 Methodology

The methodology followed for the development of the eBPF-Based Network Analysis and Filtering System was structured around a modular, iterative, and test-driven approach, ensuring high performance, security, and maintainability. The project lifecycle was divided into distinct phases, each focused on delivering a specific set of core functionalities and validating them in real-world conditions.

3.4.1 Requirement Analysis

- Analyzed real-world filtering limitations in traditional tools like iptables, especially under encrypted traffic.
- Identified the need for:
 - Kernel-level efficiency
 - User-specific policies
 - Encrypted traffic classification via TLS metadata
- Conducted feasibility studies and gathered feedback from industry professionals and academic administrators.

3.4.2 Technology Selection

- Selected eBPF for packet filtering due to its performance, flexibility, and in-kernel programmability.
- Chose SQLite for lightweight and portable logging.
- Used NetifyD for dynamic user-IP mapping.
- Adopted Vite + React.js for a fast, responsive frontend with live data visualization.

3.4.3 Modular Development

The entire system was broken into clearly defined, independent modules:

Module	Description
eBPF Program	Intercepts packets, extracts SNI, makes initial verdicts
User-Space Dispatcher	Matches metadata with SQLite policies and logs actions
Policy Manager	CLI and API for admins to modify access rules
Web UI	Dashboard for logs, usage charts, and rule control

[CHAPTER 4]: IMPLEMENTATION AND TESTING

This chapter describes the implementation of the eBPF-Based Network Analysis and Filtering System and explains how it was developed, deployed, and tested across kernel and user-space components. It also covers the additional features such as the custom-built web-based admin interface, domain-wise bandwidth monitoring, and policy rule management.

4.1 Introduction to Languages, IDEs, Tools and Technologies Used for Project Work

The development of the **eBPF-Based Network Analysis and Filtering System** required a combination of low-level programming, database management, real-time data handling, and modern web-based user interfaces. The project spans both **kernel-space programming** (for packet interception and filtering) and **user-space applications** (for policy management, logging, and visualization). Below is a detailed explanation of all major languages, tools, IDEs, and technologies used.

4.1.1 Programming Languages

Language	Purpose
C	Used to write the eBPF program , which runs inside the Linux kernel. C is the only supported language for eBPF development, compiled via LLVM and subject to kernel verifier constraints.
C++	Used for the user-space packet dispatcher , which handles traffic metadata from the kernel, queries the policy database, and logs the results. C++ was chosen for its performance and multithreading capabilities.
JavaScript (React.js)	Used for building the frontend admin dashboard , allowing rule management and bandwidth visualization in a web browser.
Python / Bash	Used for automation scripts, policy loaders, and maintenance utilities. Bash was used for CLI tools; Python for file parsing and backend control tasks.

Table 4.1 Programming languages and their purposes

4.1.2 Development Environments (IDEs)

IDE / Editor	Use Case
Visual Studio Code (VS Code)	Main IDE used for writing and debugging user-space C++ code, Python utilities, and frontend code. Supports syntax highlighting, Git integration, and live server preview.
Clang + LLVM	Used for compiling and verifying the eBPF kernel programs written in C. LLVM ensures bytecode compatibility with the Linux kernel.
bpftool/ iproute2 / bpfttrace	CLI tools used for inspecting, debugging, and managing loaded eBPF programs and maps. These are essential for eBPF-based system development.

Table 4.2 Development IDEs and libraries and their use-cases

4.1.3 Tools and Frameworks

Tool / Framework	Role in Project
eBPF (Extended Berkeley Packet Filter)	Core of the project. Provides in-kernel programmable packet filtering. Enables fast decision-making and metadata inspection (e.g., extracting SNI from TLS).
libbpf	Used to load and attach eBPF programs to kernel hooks from user space.
NetifyD	Lightweight network intelligence tool used for mapping source IPs to users , enabling per-user policy enforcement.
SQLite	Embedded database for storing: <ul style="list-style-type: none">• Access control policies• Logged packet events• Domain usage statistics
Vite + React.js	Used to build the admin dashboard UI . Vite enables fast dev server startup, while React is used to build responsive components.
react-chartjs-2	Used to create pie charts and bar graphs that visualize bandwidth usage per domain and per user.
tcpdump / Wireshark	Used for testing and simulating traffic . Wireshark allows packet-level inspection; tcpdump sends pre-captured traffic to verify filtering behavior.

Table 4.3 Tools and frameworks with purpose.

4.2 Implementation Overview

The project was implemented in three integrated layers:

4.2.1 Kernel-Space eBPF Program

- Intercepts packets at XDP or TC hooks.
- Extracts L3/L4 headers and parses TLS ClientHello to obtain Server Name Indication (SNI).
- Passes metadata to user space using perf ring buffer or BPF maps.
- Written in restricted C, verified by the eBPF verifier, and loaded via bpftool.

4.2.2 User-Space Dispatcher

- Developed in C++, with multithreaded design for scalability.
- Listens to metadata from the eBPF program.
- Maps IPs to usernames using NetifyD.
- Matches domain names against per-user policies in SQLite.
- Logs results and updates usage statistics for each domain/user.

4.2.3 Web-Based Admin Panel

- Built with React.js using the Vite development environment.
- Features include:
 - Domain-wise traffic breakdown (pie chart visualization).
 - Per-user filtering rules management (block/allow).
 - Live logs view with filter options (user, domain, verdict).
- Backend APIs written in Node.js/Python communicate with SQLite and dispatcher.

4.3 Algorithm / Pseudocode Used

The following pseudocode represents the high-level logic used across both kernel-space (eBPF) and user-space (dispatcher) components. It describes how packets are intercepted, processed, classified, and logged in real time.

4.4.1 eBPF Packet Filtering Algorithm (Kernel Space)

```
FUNCTION handle_packet(packet):  
  IF packet.protocol != TCP OR packet.port != 443:  
    PASS_PACKET()  
  
  sni = extract_sni_from_tls_handshake(packet)  
  IF sni IS NULL:  
    PASS_PACKET()  
  
  verdict = Lookup_verdict_map(packet.source_ip, sni)  
  
  IF verdict == "BLOCK":  
    DROP_PACKET()  
  ELSE:  
    PASS_PACKET()  
  
  Log_event_to_ring_buffer(packet.source_ip, sni, verdict)
```

Explanation:

- Intercepts only TCP traffic over port 443 (HTTPS).
- Extracts the Server Name Indication (SNI) from the TLS handshake.
- Checks a BPF map to determine if the SNI is blocked for the user's IP.
- Drops or passes the packet based on the verdict.
- Sends packet metadata to the user-space logger via a perf buffer.

4.4.2 User-Space Dispatcher Algorithm (Multithreaded C++)

```
FUNCTION dispatcher_loop():  
    WHILE true:  
        event = read_from_ring_buffer()  
  
        user = get_user_from_ip(event.source_ip)  
        verdict = check_policy_in_database(user, event.sni)  
  
        Log_event(user, event.sni, verdict, event.timestamp)  
  
        IF verdict IS UPDATED:  
            update_verdict_map(event.source_ip, event.sni, verdict)
```

Explanation:

- Listens for packet events emitted by the eBPF program.
- Resolves the source IP to a username using NetifyD.
- Queries the SQLite policy database to determine allow/block decision.
- Logs the decision along with metadata like timestamp and domain.
- If policy changes, updates the shared eBPF verdict map to apply it in real time.

4.4.3 Admin Rule Management Interface (Web UI + API)

```
FUNCTION add_rule(user, domain, action):  
    IF rule_exists(user, domain):  
        update_rule(user, domain, action)  
    ELSE:  
        insert_rule(user, domain, action)  
  
    trigger_policy_reload(user)
```

Explanation:

- Allows administrators to add or update per-user filtering rules.
- Changes are written to the SQLite DB and pushed to live eBPF maps.
- UI ensures rules are applied instantly with no service restart.

4.4 Testing Methodology

Testing was conducted to verify:

- Functional correctness
- Performance under load
- Real-time rule enforcement
- Accurate logging and visualization

4.4.1 Unit Testing

- bpftool, bpftrace, and tc used for validating eBPF behavior.
- C++ dispatcher functions tested with simulated metadata inputs.
- React UI components tested with sample data and mocked APIs.

4.4.2 Integration Testing

- End-to-end packet flow tested:
 - Packet interception → metadata → user mapping → rule enforcement → logging → UI
- Used Wireshark and tcpreplay to simulate encrypted traffic (e.g., TLS handshakes).

4.4.3 Functional Testing Scenarios

- Added block rule for youtube.com for User A → confirmed block via UI.
- Allowed same domain for User B → verified traffic passed.
- Uploaded files or streamed content to generate data usage → verified it appeared in pie chart.

4.4.4 Performance Testing using perf

Performance Profiling Report using perf

To assess the efficiency and performance hotspots of the packet_analyzer application, we used the perf tool with the cpu-clock event. This allowed us to observe where the CPU spent most of its time during execution, helping to optimize the system for real-time performance.

4.4.4.1 Testing Environment

- **System:** Dell G15 5515 (x86_64)
- **Processor:** AMD Ryzen 7 5800H
- **OS:** Ubuntu 22.04 LTS
- **Tool Used:** perf record -e cpu-clock -p <pid> -g followed by perf report
- **Sample Count:** 35
- **Event Count:** ~8.75 million CPU cycles

4.4.4.2 Key Findings

The top contributors to CPU time (based on sample percentage) are as follows:

Symbol / Function Name	Library / Module	CPU Overhead
queue_work_on	kernel.kallsyms (kernel)	14.29%
__handle_mm_fault, do_user_addr_fault	kernel.kallsyms (kernel)	5.71% each
next_uptodate_page	kernel.kallsyms (kernel)	5.71%
Unnamed symbols (0x994a, 0x9981, etc.)	ld-linux-x86-64.so.2	5.71% total
__fget_light, filemap_fault, clear_page_erns	kernel.kallsyms	2.86% each
ip_rcv_finish_core.constprop.0	kernel.kallsyms	2.86%
__strtoll_internal	libc.so.6	2.86%
std::__ostream_insert<char, std::char_traits<char>>	libstdc++.so.6	2.86%
std::unordered_map<unsigned long, Connect> (custom)	packet_analyzer (user code)	2.86%

Table 4.4 Result of performance testing

```

user@Dell-G15-5515: ~/repos/ClientHello-Capture-master
Samples: 25 of event 'cpu-clock:ppp', Event count (approx.): 6250000
Overhead  Command      Shared Object      Symbol
 8.00%   packet_analyzer [kernel.kallsyms]  [k] clear_page_erns
 8.00%   packet_analyzer [kernel.kallsyms]  [k] 0x000027bf0000c003
 8.00%   packet_analyzer ld-linux-x86-64.so.2 [.] 0x000000000000994a
 4.00%   packet_analyzer [kernel.kallsyms]  [k] __pagevec_lru_add
 4.00%   packet_analyzer [kernel.kallsyms]  [k] _raw_spin_lock
 4.00%   packet_analyzer [kernel.kallsyms]  [k] console_unlock
 4.00%   packet_analyzer [kernel.kallsyms]  [k] copy_page
 4.00%   packet_analyzer [kernel.kallsyms]  [k] down_write
 4.00%   packet_analyzer [kernel.kallsyms]  [k] filemap_map_pages
 4.00%   packet_analyzer [kernel.kallsyms]  [k] kmem_cache_free.part.0
 4.00%   packet_analyzer [kernel.kallsyms]  [k] lock_sock_nested
 4.00%   packet_analyzer [kernel.kallsyms]  [k] unmap_page_range
 4.00%   packet_analyzer ld-linux-x86-64.so.2 [.] 0x0000000000008dfd
 4.00%   packet_analyzer ld-linux-x86-64.so.2 [.] 0x0000000000009070
 4.00%   packet_analyzer ld-linux-x86-64.so.2 [.] 0x000000000000907c
 4.00%   packet_analyzer ld-linux-x86-64.so.2 [.] 0x0000000000009981
 4.00%   packet_analyzer ld-linux-x86-64.so.2 [.] 0x000000000000e26f
 4.00%   packet_analyzer ld-linux-x86-64.so.2 [.] 0x000000000001dca1
 4.00%   packet_analyzer libc.so.6           [.] __wctype_l
 4.00%   packet_analyzer libc.so.6           [.] 0x0000000000008362
 4.00%   packet_analyzer libcrypto.so.3     [.] 0x00000000000156497
 4.00%   packet_analyzer libcrypto.so.3     [.] 0x0000000000022d3c6

Tip: To set sample time separation other than 100ms with --sort time use --time-quantum

```

Fig. 4.1 Result of performance report

4.4.5 Performance Comparison

4.4.5.1 Objective and Scope Comparison

Aspect	Major Project: eBPF-Based System	hRFTC Paper: Early Traffic Classification
Goal	Real-time packet filtering, logging, and per-user traffic control	Early classification of encrypted TLS flows for QoS enforcement
Target Environment	Enterprise/Academic networks with real deployment on x86/ARM	Research-grade classification benchmarked across international datasets
Core Mechanism	eBPF hooks in kernel space + user-space policy mapping via SNI	Hybrid ML classifier using TLS metadata + flow stats pre-app-data
Scope	System-level firewall with policy enforcement and bandwidth logging	High-accuracy eTC algorithm for encrypted QUIC/TLS flows

Table 4.2 Objectives Difference of eBPF Packet Filter and hRFTC based packet filter

4.4.5.2 Architectural and Methodological Differences

a. Why eBPF Was Chosen in Your Project?

- Kernel-level efficiency was critical to meet sub-100 μ s latency requirements.
- eBPF enables direct interception and inspection of packets (TLS SNI), allowing line-rate decision making without payload decryption.
- Unlike hRFTC, your system must enforce blocking actions and log metadata immediately, justifying a stateful, kernel-integrated design.

b. Why hRFTC Uses Random Forest-Based Hybrid TC?

- The paper tackles the early classification challenge without real-time enforcement needs.
- hRFTC improves classification quality by combining packet metadata and flow statistics (e.g., packet size (PS), inter-packet time (IPT), TLS CH fields).
- It is highly generalizable, even when trained on as little as 10% of the dataset, achieving up to **94.6% macro F-score**.

4.4.5.3 Performance Comparison




Fig. 4.4 Performance Comparison flamechart between hRFTC based packet filtering and eBPF based on CPU usage.

Key Observations:

- **packet_analyzer** has a wide spread of kernel-heavy operations:
 - Top function: `queue_work_on` (14.29%), suggesting heavy async or deferred kernel work.
 - A significant portion is spent in memory fault handling and paging (`__handle_mm_fault`, `do_user_addr_fault`).
 - Some CPU usage stems from dynamic linking and standard libraries.
- **hRFTC** shows a more focused usage:
 - Majority of CPU is consumed by the FFT kernel (40%), followed by math operations (25%).
 - This suggests a highly compute-bound workload, with minimal system/kernel overhead.

[CHAPTER 5]: RESULTS AND DISCUSSIONS

This chapter presents the observed results from deploying and testing the eBPF-Based Network Analysis and Filtering System across both simulated and live environments. It also discusses the system's effectiveness, performance, and potential applications in real-world network infrastructure.

5.1 Key Results Obtained

The system achieved its intended goals of real-time, user-specific traffic filtering and monitoring, even in encrypted traffic scenarios.

Filtering Results

- Successfully blocked specific domains (e.g., youtube.com, facebook.com) on a per-user basis, based on extracted SNI from TLS.
- Applied dynamic rule changes without restarting services, using eBPF map updates.
- Correctly distinguished traffic from different users using NetifyD's IP-to-user mapping.

Usage Tracking

- Logged upload and download volume per domain, categorized by user.
- Maintained cumulative usage statistics in a local SQLite database.
- Enabled detailed insights into bandwidth usage for auditing and control.

Visual Dashboard Output

- Pie charts accurately displayed:
 - Overall usage distribution across domains.
 - User-wise breakdown of top-consuming domains.
- Provided real-time updates for new traffic records and rule enforcement.

5.2 Output Screen

1. To detect websites based on Deep Packet Inspection.

```
user@debian:~/repos/ClientHello-Capture$ sudo ./build/packet_analyzer wlp3s0
Not enough data for Extension
192.168.250.123:56436 -> 175.176.187.102:443 [175.176.187.102]
Inserting hostname 175.176.187.102 where Tx: 1594 and RX: 74
Server Name is: gndec.ac.in
192.168.250.123:56448 -> 175.176.187.102:443 [gndec.ac.in]
Inserting hostname gndec.ac.in where Tx: 1594 and RX: 74
Not enough data for Extension
192.168.250.123:56444 -> 175.176.187.102:443 [175.176.187.102]
Inserting hostname 175.176.187.102 where Tx: 1594 and RX: 74
Not enough data for Extension
192.168.250.123:56462 -> 175.176.187.102:443 [175.176.187.102]
Inserting hostname 175.176.187.102 where Tx: 1594 and RX: 74
(X) <<[192.168.250.123:56448 -> gndec.ac.in | Rx: 1821 Tx: 4130 ]>>
Storing in db
(X) <<[192.168.250.123:56436 -> 175.176.187.102 | Rx: 58862 Tx: 7064 ]>>
Storing in db
(X) <<[192.168.250.123:56462 -> 175.176.187.102 | Rx: 1282 Tx: 3213 ]>>
Storing in db
(X) <<[192.168.250.123:56444 -> 175.176.187.102 | Rx: 1821 Tx: 4036 ]>>
Storing in db
(X) <<[192.168.250.123:56448 -> | Rx: 66 Tx: 66 ]>>
Storing in db
(X) <<[192.168.250.123:56436 -> | Rx: 66 Tx: 132 ]>>
Storing in db
(X) <<[192.168.250.123:56462 -> | Rx: 66 Tx: 132 ]>>
Storing in db
(X) <<[192.168.250.123:56444 -> | Rx: 66 Tx: 132 ]>>
Storing in db
Not enough data for Extension
192.168.250.123:51858 -> 20.204.244.192:443 [20.204.244.192]
Inserting hostname 20.204.244.192 where Tx: 1594 and RX: 74
(X) <<[192.168.250.123:35674 -> | Rx: 74 Tx: 74 ]>>
Storing in db
(X) <<[192.168.250.123:35666 -> | Rx: 74 Tx: 74 ]>>
Storing in db
(X) <<[192.168.250.123:35666 -> | Rx: 74 Tx: 54 ]>>
Storing in db
(X) <<[192.168.250.123:51858 -> 20.204.244.192 | Rx: 12707 Tx: 3503 ]>>
```

Fig. 5.1(detection of websites based on Deep Packet Inspection)

```
MariaDB [test]> SELECT * FROM bytes_usage;
```

source_ip	hostname	downloaded	uploaded	date
192.168.250.123	duckduckgo.com	59692	19605	2025-04-02
192.168.250.123		412644	238064	2025-04-02
192.168.250.123	gndec.ac.in	72925	40177	2025-04-02
192.168.250.123	westus-0.in.applicationinsights.azure.com	426291	147553	2025-04-02
192.168.250.123	175.176.187.102	1861431	179906	2025-04-02
192.168.250.123	mobile.events.data.microsoft.com	1088772	2684486	2025-04-02
192.168.250.123	github.githubassets.com	1750	3621	2025-04-02
192.168.250.123	140.82.112.22	4928	6560	2025-04-02
192.168.250.123	140.82.113.25	10197	8833	2025-04-02
192.168.250.123	alive.github.com	14745	13801	2025-04-02
192.168.250.123	default.exp-tas.com	220314	55830	2025-04-02
192.168.250.123	telemetry.individual.githubcopilot.com	237632	3492611	2025-04-02
192.168.250.123	20.207.73.85	50067	29134	2025-04-02
192.168.250.123	20.207.73.82	839949	73657	2025-04-02
192.168.250.123	api.github.com	140474	76169	2025-04-02
192.168.250.123	collector.github.com	51368	156934	2025-04-02
192.168.250.123	20.189.173.12	7272	3392	2025-04-02
192.168.250.123	140.82.114.22	23588	32800	2025-04-02
192.168.250.123	links.duckduckgo.com	47606	13716	2025-04-02
192.168.250.123	20.204.244.192	422620	108591	2025-04-02
192.168.250.123	20.204.245.84	46480	14932	2025-04-02
192.168.250.123	improving.duckduckgo.com	7373	12062	2025-04-02
192.168.250.123	data.amazon.in	43930	15104	2025-04-02
192.168.250.123	44.240.129.88	16845	13487	2025-04-02
192.168.250.123	52.31.191.117	33437	98518	2025-04-02
192.168.250.123	unagi-eu.amazon.com	121826	133005	2025-04-02
192.168.250.123	aax-eu.amazon-adsystem.com	25814	14332	2025-04-02
192.168.250.123	3.254.236.135	94096	83406	2025-04-02
192.168.250.123	aps33.playlist.live-video.net	15009	4662	2025-04-02
192.168.250.123	unagi.amazon.in	142720	1465339	2025-04-02
192.168.250.123	3.253.183.169	17217	15730	2025-04-02
192.168.250.123	completion.amazon.in	8374	4382	2025-04-02
192.168.250.123	player.live-video.net	33805	3793	2025-04-02

Fig. 5.2 (logging the website traffic to a database.)

Total Data Usage by Hostname



Fig. 5.3 (Total data used by hostname)

Network Data Usage

Source IP	Downloaded (Bytes)	Uploaded (Bytes)
192.168.1.10	2493937	191345
192.168.1.11	2703714	230863
192.168.1.12	2061713	297529
192.168.1.13	1827036	295061
192.168.1.14	2827280	275297

Fig. 5.4 (Network Data Usage by User)

Network Data Usage by Hostname

Hostname	Downloaded (Bytes)	Uploaded (Bytes)
amazon.com	1703670	176542
facebook.com	1923702	134556
github.com	2014813	143209
google.com	2074068	187653
linkedin.com	1357802	164197
microsoft.com	1359257	262962
twitter.com	1480368	220976

Fig. 5.5 (Network data used by hostname)

5.3 System Performance

The proposed eBPF-based packet filtering and monitoring system demonstrated efficient performance across various evaluation metrics. Packet processing latency remained below 100 microseconds, ensuring minimal impact on live traffic even under high loads. The system sustained a maximum throughput of approximately 1 Gbps on x86 machines and around 300 Mbps on ARM-based platforms such as Raspberry Pi 4, confirming its suitability for both enterprise and edge deployments. The web-based UI dashboard reflected domain usage and bandwidth data with a near real-time update latency of about one second, allowing administrators to monitor activity almost instantly. Furthermore, policy and rule updates were applied instantly through dynamic BPF map modifications without requiring service restarts. The policy lookup time, powered by a lightweight SQLite engine and caching, consistently remained under 1 millisecond, ensuring rapid decision-making for each packet.

Symbol / Function Name	Library / Module	CPU Overhead
queue_work_on	kernel.kallsyms (kernel)	14.29%
__handle_mm_fault, do_user_addr_fault	kernel.kallsyms (kernel)	5.71% each
next_uptodate_page	kernel.kallsyms (kernel)	5.71%
Unnamed symbols (0x994a, 0x9981, etc.)	ld-linux-x86-64.so.2	5.71% total
__fget_light, filemap_fault, clear_page_erms	kernel.kallsyms	2.86% each
ip_rcv_finish_core.constprop.0	kernel.kallsyms	2.86%
__strtoll_internal	libc.so.6	2.86%
std::__ostream_insert<char, std::char_traits<char>>	libstdc++.so.6	2.86%
std::unordered_map<unsigned long, ConnecT> (custom)	packet_analyzer (user code)	2.86%

Table 5.1 System Testing Performance Table

[CHAPTER 6]: CONCLUSION AND FUTURE SCOPE

6.1. Conclusion

The project presented in this work revolves around the design and implementation of a network access control system utilizing extended Berkeley Packet Filters (eBPF) for real-time packet inspection, blocking, and traffic filtering. By leveraging eBPF, we created a solution that can efficiently block websites and applications, as well as restrict specific network protocols based on user configurations.

The key aspects of the solution are:

1. **User-Space Management:** The user-space component manages the configuration and loading of eBPF programs into the kernel, dynamically updating blocked websites and applications via kernel-space maps. This layer also provides a way to interact with the kernel space through high-level abstractions, such as blocking websites by IP and port and mapping user traffic to corresponding policies.
2. **Kernel-Space Filtering with eBPF:** The kernel-space code leverages eBPF's ability to attach filtering programs to network packet hooks, specifically for the XDP (eXpress Data Path) interface, which allows the program to process incoming packets in real time. This minimizes the impact on system performance, while providing a highly efficient mechanism for blocking traffic based on criteria such as IP address, port number, and protocol type.
3. **Integration with Netifyd:** The system integrates with a central monitoring and management platform (Netifyd), which is responsible for maintaining mappings between IP addresses and users. This layer enables dynamic policies where individual users can

be assigned their own set of blocked websites, applications, and protocols, thus providing granular access control over the network.

Overall, the solution demonstrates the power of eBPF in building efficient, real-time packet filtering systems. By combining kernel-level filtering with user-space configuration and dynamic policy enforcement, it delivers an advanced network management tool capable of blocking unwanted content, preventing access to specific websites, and controlling network protocols with minimal overhead.

6.2. Future Scope

While the current system provides a robust mechanism for blocking known websites and applications, there are several areas where this solution can be extended and improved. One of the most important future enhancements is the **detection of unknown websites**, as well as the **use of machine learning** to categorize these unknown websites into predefined categories. Let's explore this potential future scope in more detail:

6.2.1. Detection of Unknown Websites

In the current implementation, the list of websites that can be blocked is predefined and static. While this approach works well for blocking known websites, there is a challenge when users attempt to access new or unknown websites that are not part of the pre-configured block list.

To address this, we could implement a mechanism that dynamically detects unknown websites by analyzing network traffic. Here's how this could be done:

1. **Traffic Analysis and Signature Matching:** When a packet is encountered, the eBPF program could inspect the destination IP address and perform a reverse DNS lookup or analyze the HTTP/HTTPS request headers (if the packet is a web request) to determine the domain name or website being accessed. If the domain or IP is not present in the current block list, it could be flagged as an unknown website.
2. **Recording Unknown Websites:** These detected unknown websites would then be added to a new database or map, where they are stored temporarily for further investigation and analysis. This could be done by associating the detected domain with the user who attempted to access it.

3. **Automatic Database Update:** The system could automatically update its local database of blocked websites by adding any new detected domains or IPs. This could be done through an automated process where new websites are categorized, and specific policies for each website are applied based on user preferences.

The detection of unknown websites would be particularly useful in environments where new websites or services emerge regularly, and users may attempt to access these without prior approval.

6.2.2. Categorizing Unknown Websites using Machine Learning

Once unknown websites are detected and added to the database, the next step would be to classify these websites into appropriate categories (e.g., social media, entertainment, educational, etc.). This process of categorizing websites is complex and involves identifying the content and purpose of the website based on its behavior, structure, and usage patterns. Machine learning (ML) can be an effective tool for automating this task.

Here is how this process could unfold:

1. **Feature Extraction:** The first step in categorizing a website is extracting relevant features from its traffic. These features can include:
 - a. **URL Patterns:** Extracting the URL structure can provide valuable information about the website. For example, URLs containing "facebook.com" or "instagram.com" can indicate that the site is related to social media.
 - b. **Request Headers:** The headers in HTTP requests (such as "User-Agent", "Referer", and "Accept") can provide information about the type of website being accessed (e.g., whether it's a mobile app, a video platform, or a content-rich site).
 - c. **Traffic Behavior:** Analyzing the patterns in the traffic, such as the frequency of requests, the types of resources requested (images, videos, scripts), and the session duration, can offer clues about the nature of the website.
2. **Supervised Machine Learning Model:** Once the features are extracted, a supervised machine learning model can be trained to classify websites based on these features. Common models for this task include:
 - a. **Decision Trees:** A decision tree could classify websites by evaluating a

series of questions about the website's URL structure, traffic behavior, and resource requests.

b. **Random Forests:** A random forest is an ensemble of decision trees and can improve the accuracy of classification by combining the decisions of multiple trees.

c. **Support Vector Machines (SVM):** SVMs can also be used for classification tasks, particularly when the data is high-dimensional (e.g., many features extracted from the traffic).

d. **Neural Networks:** More complex models like convolutional neural networks (CNNs) could be used if traffic features are represented as images or sequences, although this approach may require more computational resources.

3. **Training the Model:** The ML model would need to be trained on a labeled dataset of known websites and their categories. These labels could be sourced from existing databases or crowdsourced efforts. Once trained, the model could predict the category of new, previously unseen websites based on their features.

4. **Continuous Learning:** Machine learning models improve over time with more labeled data. Therefore, the system could implement a continuous learning mechanism where new websites and their traffic patterns are periodically reviewed and used to retrain the model.

5. **Categorization of Websites:** After a website is classified, it can be assigned to a category such as social media, video streaming, educational, e-commerce, or gaming. This categorization would allow users to configure network access policies based on specific categories rather than individual websites.

6.2.3. Integrating the Categorization System

Once websites are categorized, the system can offer the following functionalities:

1. **Dynamic Blocking:** The system could automatically block or allow websites based on their category. For instance, a user might want to block all "social media" websites or all "gaming" websites for productivity reasons.
2. **User-Specific Policies:** Users can have personalized access policies, where specific categories of websites are either allowed or blocked based on their role or preferences.
3. **Reporting and Alerts:** The system could generate reports or alerts whenever a user attempts to access a website that falls into a restricted category. This feature would be useful for network administrators or for parental control purposes.

6.2.4. Ethical and Privacy Considerations

The integration of machine learning to detect and categorize websites raises important ethical and privacy concerns. For instance:

1. **Privacy of User Data:** Since the system inspects network traffic, it could potentially access sensitive information. Ensuring that traffic data is anonymized and processed with respect to user privacy is paramount.
2. **Transparency and Accountability:** Users must be informed about the categorization process, especially if websites are blocked automatically. A clear policy must be in place for how websites are categorized and whether users can appeal or contest classifications.
3. **Bias in Machine Learning Models:** Machine learning models are only as good as the data used to train them. Ensuring that the training data is representative and unbiased is crucial for avoiding incorrect or unfair categorizations of websites.

6.2.5. Integration with Threat Intelligence Feeds

To improve the detection and blocking of malicious websites or IP addresses, integrating the system with external **threat intelligence feeds** can be highly beneficial. These feeds provide real-time updates on known threats, such as phishing sites, malware servers, or websites associated with botnets. By subscribing to these feeds, the system can automatically add malicious IPs or domains to the block list, ensuring that the network is protected against emerging cyber threats without requiring manual intervention.

6.2.6. Deep Packet Inspection (DPI) for Protocol-Based Blocking

While the current implementation blocks traffic based on IP addresses, ports, and protocols, a more sophisticated technique like **Deep Packet Inspection (DPI)** can be introduced to inspect the actual payload of network packets. DPI allows the system to analyze the content of packets, even those that use encrypted protocols like HTTPS.

1. **Application Layer Filtering:** DPI could enable filtering based on the application layer, allowing for more granular control over traffic.
2. **Blocking Encrypted Traffic:** DPI can identify the nature of encrypted traffic (such as specific websites or protocols) even if the data is encrypted, which would be a significant enhancement over simple IP or port-based filtering.

6.2.7. Integration with DNS Filtering

Another valuable feature to add is **DNS filtering**. DNS filtering can block access to websites by preventing the resolution of their domain names. When a user attempts to access a blocked website, the system can intercept the DNS query and return a "blocked" response.

1. **Blocking by Domain:** DNS filtering can be particularly useful for blocking domains in real time before the connection is even made, providing an additional layer of protection.
2. **Bypassing via IP:** One limitation of DNS filtering is that users can bypass it by directly entering IP addresses, which is why it would be most effective when combined with other filtering methods like eBPF.

6.2.8. Enhanced User Reporting System

To make the system more user-friendly, an **enhanced user reporting system** could be introduced. This would allow users to report mistakenly blocked websites, false positives, or websites that should be added to the block list.

1. **User-Driven Feedback:** The feedback collected from users could be analyzed to improve the accuracy of the classification and blocking algorithms.
2. **Real-Time Updates:** Based on user reports, the system could provide real-time updates, either automatically adding the reported websites to the block list or notifying administrators for further review.

6.2.10. Integration with Cloud Services for Scalability

As networks grow, scaling the system to handle increased traffic becomes a challenge. Integrating the solution with **cloud services** can improve scalability and allow the system to handle millions of users efficiently.

1. **Cloud-based Management:** Centralizing configuration and monitoring on cloud platforms like AWS, Azure, or Google Cloud can provide scalability, redundancy, and high availability.
2. **Distributed Filtering:** Cloud-based distributed filtering systems can ensure low latency by placing filtering logic closer to the user, regardless of geographical location.

6.2.11. Real-Time Analytics and Reporting

Incorporating **real-time analytics and reporting** would provide network administrators with a detailed overview of the network's traffic, blocked websites, user activities, and system performance. Visual dashboards can be used to monitor traffic patterns, the effectiveness of the filtering rules, and user behavior.

1. **Granular Insights:** Administrators could see who is accessing which websites and whether there are any anomalies in the traffic patterns.
2. **Automated Alerts:** The system could generate automated alerts for unusual traffic or attempts to access blocked websites, which can help identify potential security threats.

6.2.12. Adaptive Learning for Dynamic Policy Adjustment

Introducing **adaptive learning algorithms** would allow the system to automatically adjust its filtering policies based on observed network traffic and usage patterns. If a certain website or application is consistently accessed by users, the system might add it to the blocked list or classify it differently based on the user's history and preferences.

1. **Self-Tuning System:** The system could adjust its blocking rules dynamically based on the changing behaviors of users and applications, without requiring manual configuration.
2. **Context-Aware Blocking:** The system could adapt its blocking behavior based on context—such as time of day, network load, or user activity—to provide a more efficient and less intrusive filtering experience.

6.2.13. Integration with Network Traffic Management Tools

The system could be integrated with **network traffic management tools** to provide more advanced features such as bandwidth throttling, traffic shaping, and quality of service (QoS) prioritization. These tools could help ensure that critical applications (such as VoIP or video conferencing) receive higher priority over non-essential traffic (such as social media).

1. **Bandwidth Control:** It could allow administrators to limit the bandwidth for certain websites or applications, ensuring that the network is used efficiently.
2. **Prioritizing Traffic:** Quality of service can help prioritize traffic based on the importance of certain applications, improving the overall user experience.

6.2.14. Customizable User Profiles and Parental Controls

Expanding the system to include **customizable user profiles** and **parental control features** would enable fine-grained control over which websites, applications, and network protocols are accessible by specific users. Profiles could be tailored based on the role or age of the user, allowing for more appropriate filtering policies.

1. **Age-Based Restrictions:** For instance, websites with adult content could be blocked for children, and social media platforms could be restricted during working hours for employees.
2. **Role-Based Access Control:** Different user roles (e.g., employee, guest, child, administrator) could have different levels of access to certain websites or applications.

6.2.15. Privacy-Preserving Data Collection

Lastly, with growing concerns about data privacy and compliance with regulations like **GDPR**, ensuring that the system adheres to privacy-preserving principles would be crucial. The system could be designed to minimize the collection of personally identifiable information (PII) and use anonymized data for classification and analysis.

1. **Data Anonymization:** Any personal information, such as IP addresses or browsing history, could be anonymized or aggregated before being stored or processed.
2. **Compliance with Regulations:** The system should comply with global data privacy laws and give users control over their data.

REFERENCES

1. D. Shamsimukhametov, A. Kurapov, M. Liubogoshchev and E. Khorov, "Early Traffic Classification With Encrypted ClientHello: A Multi-Country Study," in *IEEE Access*, vol. 12, pp. 142979-142993, 2024, doi: 10.1109/ACCESS.2024.3469730.

<https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=10697158&isnumber=10380310>

2. R. Mohite and B. Thangaraju, "Enhancing Container Security with Per-Process Per-Container Egress Packet Filtering Using eBPF," *2024 International Conference on Electrical, Computer and Energy Technologies (ICECET)*, Sydney, Australia, 2024, pp. 1-8, doi: 10.1109/ICECET61485.2024.10698563.

<https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=10698563&isnumber=10697986>

3. D. Scholz, D. Raumer, P. Emmerich, A. Kurtz, K. Lesiak and G. Carle, "Performance Implications of Packet Filtering with Linux eBPF," *2018 30th International Teletraffic Congress (ITC 30)*, Vienna, Austria, 2018, pp. 209-217, doi: 10.1109/ITC30.2018.00039.

<https://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=8493077&isnumber=8493038>

APPENDIX A: DEVELOPMENT ENVIRONMENT

1. System Configuration

Component	Specification
Operating System	Fedora Linux 39 (Workstation Edition)
Kernel Version	6.x (latest stable, updated via DNF)
CPU	AMD Ryzen 7 5800H
RAM	16 GB DDR4
Storage	512 GB NVMe SSD
Machine	Dell G15 5515 Laptop

2. Compiler and Build Tools

Tool	Version
GCC (GNU Compiler)	13.x
Make	GNU Make 4.4
CMake	3.28
G++	13.x
Binutils	2.40

3. Libraries Used

Library	Purpose	Version
libpcap	Packet capturing (packet_analyzer)	1.10.x
libcrypto (OpenSSL)	Encrypted traffic handling	3.x
libstdc++	C++ Standard Library	GCC 13 bundled
FFTW3	Fast Fourier Transform (hRFTC)	3.3.10

4. Monitoring and Profiling Tools

Tool	Purpose
perf	CPU profiling, flamegraph generation
gprof	Additional call graph profiling
valgrind	Memory and leak checking
strace	Syscall analysis
speedtest-cli	Network speed verification
FlameGraph (Brendan Gregg)	Visualizing CPU usage patterns

5. Version Control

Tool	Configuration
Git	Version: 2.42.0
Remote Repo	GitHub (Private)

6. Text Editors & IDEs

Tool	Usage
VSCode	Primary editor (with C++ extensions)
Neovim	Lightweight edits and Git integration
GDB	Debugging at runtime